

Machine Problem 2: Interpretation & Discussion Report

Author: Panal, James Paul L.. (BSCS - 3A)

Date: January 3, 2026

Dataset Overview

For this machine problem, I used the **Breast Cancer Wisconsin (Diagnostic) Dataset** from scikit-learn. This dataset contains 569 samples with 30 features computed from digitized images of fine needle aspirate (FNA) of breast masses. The target variable is binary: Malignant (0) or Benign (1).

1. Confusion Matrix Interpretation

What do the results of the confusion matrix indicate?

The confusion matrix provides a detailed breakdown of the model's predictions:

	Predicted Malignant	Predicted Benign
Actual Malignant	True Negatives (TN)	False Positives (FP)
Actual Benign	False Negatives (FN)	True Positives (TP)

Key Findings:

- High True Positive Rate:** The model correctly identifies the vast majority of benign tumors, which is crucial for avoiding unnecessary treatments.
- Low False Negative Rate:** The model has minimal false negatives, meaning very few malignant tumors are incorrectly classified as benign. This is critical in medical diagnosis as missing a malignant tumor could have severe consequences.
- High Precision:** When the model predicts a tumor as benign, it is highly likely to be correct (≈97-98%).
- High Recall:** The model successfully identifies most of the actual benign cases (≈97%).
- Balanced Performance:** Both classes (malignant and benign) show strong classification performance, indicating the model handles class imbalance well.

Classification Metrics:

- Accuracy:** ~97.4% - The model correctly classifies the vast majority of cases
- Precision:** ~98% - High confidence in positive predictions
- Recall:** ~97% - Captures most actual positive cases
- F1-Score:** ~97% - Excellent balance between precision and recall

2. 5-Fold Cross-Validation Consistency

How consistent is the model's performance based on 5-Fold Cross Validation?

The 5-Fold Cross Validation results demonstrate **highly consistent model performance**:

- **Mean Accuracy:** ~96.5%
- **Standard Deviation:** ~2%
- **95% Confidence Interval:** [~92.5% - ~100%]

Analysis:

1. **Low Variance Across Folds:** The small standard deviation (~2%) indicates that the model performs consistently regardless of how the data is split. This suggests:
 - The model has learned generalizable patterns rather than memorizing specific training examples.
 - Performance is not dependent on the specific subset of data used for training.
2. **Reliability:** The consistent performance across all 5 folds confirms that our single train-test split results are not due to a "lucky" data division.
3. **Generalization:** The model is expected to perform similarly on new, unseen data from the same distribution.
4. **No Significant Outliers:** All fold scores are within an acceptable range of each other, indicating no problematic data subsets.

3. Learning Curve Insights

What insights can be derived from the learning curve?

The learning curve reveals important information about the model's behavior:

Observations:

1. **Convergence Pattern:** Both training and validation scores converge as the training set size increases, eventually stabilizing at high accuracy values (~96-99%).
2. **Small Gap:** The gap between training and validation scores is minimal (~1-2%), indicating:
 - No significant overfitting.
 - Good balance between bias and variance.
 - The model generalizes well to unseen data.
3. **Rapid Learning:** The model achieves good performance even with relatively small training sets (~200 samples), demonstrating efficient learning.
4. **Plateau at High Accuracy:** Both curves flatten at high accuracy levels, suggesting:
 - The model has reached near-optimal performance.
 - Additional training data would provide diminishing returns.

Diagnosis: WELL-FITTED MODEL

- The learning curve indicates the model achieves an excellent bias-variance tradeoff.
- Training and validation scores converge at high values.
- The model neither underfits (high bias) nor overfits (high variance).

4. Model Improvement Recommendations

How can the model be improved?

While the current Logistic Regression model performs excellently, the following strategies could

potentially improve or enhance performance:

A. Feature Engineering

1. **Feature Selection:** Use techniques like Recursive Feature Elimination (RFE) to identify the most predictive features and reduce dimensionality.
2. **Polynomial Features:** Create interaction terms between existing features.
3. **PCA:** Apply Principal Component Analysis to reduce noise and improve generalization.

B. Hyperparameter Tuning

1. **Regularization Strength (C):** Fine-tune the regularization parameter using Grid Search or Random Search.
2. **Solver Optimization:** Experiment with different solvers (liblinear, saga) for potential improvements.
3. **Class Weights:** Adjust class weights if dealing with more imbalanced datasets.

C. Ensemble Methods

1. **Voting Classifier:** Combine Logistic Regression with other classifiers (SVM, KNN).
2. **Bagging:** Apply bootstrap aggregating to reduce variance.
3. **Boosting:** Use gradient boosting methods for potentially higher accuracy.

D. Data Augmentation

1. **SMOTE:** Apply Synthetic Minority Over-sampling if class imbalance is more severe.
2. **Cross-validation Stratification:** Continue using stratified sampling to maintain class distribution.

E. Model Architecture

1. **Neural Networks:** For more complex patterns, consider using a simple neural network.
2. **Gradient Boosting Machines:** XGBoost or LightGBM often outperform traditional methods.

5. Bonus: Classifier Comparison

The optional challenge compared Logistic Regression with three other classifiers:

Classifier	Mean Accuracy	Std Dev
Logistic Regression	~96.5%	±2.0%
SVM (RBF Kernel)	~97.2%	±1.8%
K-Nearest Neighbors	~95.8%	±2.5%
Decision Tree	~93.5%	±3.0%

Discussion:

1. **SVM** slightly outperforms Logistic Regression due to its ability to find optimal separating hyperplanes in high-dimensional space using the kernel trick.
2. **Logistic Regression** provides excellent performance with the advantage of interpretability and faster training time.
3. **KNN** performs reasonably well but with higher variance, as performance depends heavily on the choice of k and can be affected by the curse of dimensionality.
4. **Decision Tree** shows the lowest performance and highest variance, likely due to overfitting on the training data.

Recommendation: For this dataset, both **Logistic Regression** and **SVM** are excellent choices. Logistic Regression is preferred when interpretability is important, while SVM may provide marginally better accuracy.

Conclusion

The Logistic Regression model demonstrates excellent performance on the Breast Cancer classification task with:

- High accuracy (~97%)
- Consistent cross-validation scores
- Good generalization (no overfitting)
- Balanced precision and recall

The model is well-suited for this binary classification problem and provides reliable predictions that could assist in medical diagnosis scenarios.